



Kafka-Spark-Flink -Projekat 2- Bike Washington

Big Data Systems

Mentor:

Prof. dr Dragan H. Stojanovic

Student:

Petrovic Nikola 1466

Kreiranje docker-compose fajla

```
services:
  kafka:
    image: wurstmeister/kafka:2.13-2.7.0
    depends_on:
      - zookeeper
    ports:
      - "9091:9091"
    expose:
      - "9092"
    environment:
      KAFKA_ADVERTISED_LISTENERS: INTERNAL://kafka:9092,EXTERNAL://localhost:9091
      KAFKA_LISTENER_SECURITY_PROTOCOL_MAP: INTERNAL:PLAINTEXT,EXTERNAL:PLAINTEXT
      KAFKA_LISTENERS: INTERNAL://0.0.0.0:9092,EXTERNAL://0.0.0.0:9091
      KAFKA_ZOOKEEPER_CONNECT: zookeeper:2181
      KAFKA_INTER_BROKER_LISTENER_NAME: INTERNAL
    volumes:
      - /var/run/docker.sock:/var/run/docker.sock
    networks:
      - bde

  zookeeper:
    image: wurstmeister/zookeeper:latest
    ports:
      - "2181:2181"
    networks:
      - bde
```

```
producer:
  build:
    context: .
    dockerfile: producer/Dockerfile
  depends_on:
    - kafka
  environment:
    SCRIPT: producer/producer.py
    DATA: data/2016Q1-capitalbikeshare-tripdata.csv
    KAFKA_HOST: kafka:9092
    KAFKA_TOPIC1: bikes-spark
    KAFKA_TOPIC2: bikes-flink
    KAFKA_INTERVAL: 1
  networks:
    - bde
```

```
spark-master:
  image: bde2020/spark-master:3.1.2-hadoop3.2
  container_name: spark-master
  ports:
    - "8070:8070"
    - "7077:7077"
  environment:
    - INIT_DAEMON_STEP=setup_spark
    - SPARK_MASTER_PORT=7077
    - SPARK_MASTER_WEBUI_PORT=8070
  networks:
    - bde
```

```
spark-worker-1:
  image: bde2020/spark-worker:3.1.2-hadoop3.2
  container_name: spark-worker-1
  depends_on:
    - spark-master
  ports:
    - "8071:8071"
  environment:
    - "SPARK_MASTER=spark://spark-master:7077"
    - SPARK_WORKER_WEBUI_PORT=8071
  networks:
    - bde
```

```
spark-worker-2:
  image: bde2020/spark-worker:3.1.2-hadoop3.2
  container_name: spark-worker-2
  depends_on:
    - spark-master
  ports:
    - "8072:8071"
  environment:
    - "SPARK_MASTER=spark://spark-master:7077"
    - SPARK_WORKER_WEBUI_PORT=8071
  networks:
    - bde
```

version: "3.9"

```
networks:
  bde:
    external: true
```

Kafka producer

```
Dockerfile • producer.py × COMMIT_EDITMSG
producer > producer.py > ...
1  import time
2  import os
3  import csv
4  import json
5  from datetime import datetime, timezone
6  from kafka import KafkaProducer
7
8  producer = KafkaProducer(
9      bootstrap_servers=[os.environ["KAFKA_HOST"]],
10     value_serializer=lambda v: json.dumps(v).encode("utf-8"),
11     api_version=(0, 11),
12 )
13
14 while True:
15     with open(os.environ["DATA"], "r") as file:
16         reader = csv.reader(file, delimiter=",")
17         headers = next(reader)
18         for row in reader:
19             value = {headers[i]: row[i] for i in range(len(headers))}
20             value["Duration"] = int(value["Duration"])
21             value["Start station number"] = int(value["Start station number"])
22             value["End station number"] = int(value["End station number"])
23
24             value["ts"] = int(time.time())
25             producer.send(os.environ["KAFKA_TOPIC1"], value=value)
26             producer.send(os.environ["KAFKA_TOPIC2"], value=value)
27             time.sleep(float(os.environ["KAFKA_INTERVAL"]))
28
```

Cassandra docker-compose

```
version: "3.10"

networks:
  bde:
    external: true

services:
  cassandra:
    image: cassandra:latest
    ports:
      - 9042:9042
    environment:
      - CASSANDRA_SEEDS=cassandra
      - CASSANDRA_CLUSTER_NAME=cassandra-cluster
      - CASSANDRA_DC=datacenter1
    networks:
      - bde
    mem_limit: 4g
```

kafka-spark-cassandra povezivanje

```
if __name__ == '__main__':

    parser = argparse.ArgumentParser()
    parser.add_argument("--N", type=int, help="The number of top start stations to select")
    args = parser.parse_args()

    N = args.N or 5 # Default to 5 if N is not provided

    conf = SparkConf()
    conf.setMaster("spark://spark-master:7077") # Podesavanje konfiguracije za kreiranje spark sesije.
    #conf.setMaster("local")
    conf.set("spark.driver.memory", "4g")

    #cassandra
    conf.set("spark.cassandra.connection.host", "cassandra")
    conf.set("spark.cassandra.connection.port", "9042") # Konfiguracija za povezivanje na cassandru.
    #conf.set("spark.cassandra.auth.username", "cassandra")
    #conf.set("spark.cassandra.auth.password", "cassandra")

    spark = SparkSession.builder.config(conf=conf).appName("Rides").getOrCreate() # Kreiranje spark sesije.

    # Get rid of INFO and WARN logs.
    spark.sparkContext.setLogLevel("ERROR")

    df = (
        spark.readStream.format("kafka")
        .option("kafka.bootstrap.servers", os.environ["KAFKA_HOST"])
        #.option("kafka.bootstrap.servers", "kafka:9092")
        .option("subscribe", os.environ["KAFKA_TOPIC"])
        #.option("subscribe", "bikes-spark")
        .option("startingOffsets", "latest")
        .option("groupIdPrefix", os.environ["KAFKA_CONSUMER_GROUP"])
        #.option("groupIdPrefix", "Spark-Group")
        .load()
    )
```

Citanje podataka sa kafke.

```
Dockerfile • consumer_spark.py • COMMIT_EDITMSG
consumer_spark > Dockerfile > ...
1 FROM bde2020/spark-python-template:3.1.2-hadoop3.2
2
3 ENV KAFKA_HOST=kafka:9092
4 ENV KAFKA_TOPIC=bikes-spark # Dodavanje kafka hosta, topic-a i grupe.
5 ENV KAFKA_CONSUMER_GROUP=Spark-Group
6 ENV SPARK_APPLICATION_PYTHON_LOCATION /app/consumer_spark.py # Definisanje .py aplikacije.
7 ENV SPARK_APPLICATION_ARGS "--N 10"
8 ENV SPARK_SUBMIT_ARGS --packages \ # Postavljenje parametara programa.
9 org.apache.spark:spark-streaming-kafka-0-10_2.12:3.1.2, \
10 org.apache.spark:spark-sql-kafka-0-10_2.12:3.1.2, \ # Dodavanje dependensija za povezivanje i komunikaciju sa kafkon
11 com.datastax.spark:spark-cassandra-connector_2.12:3.2.0 \ | Dependensi za povezivanje i komunikaciju sa kasandrom
12 --executor-memory 1G --executor-cores 1
13
```


Spark aplikacija

```
schema = StructType([\n    StructField("Duration", StringType(), False),\n    StructField("Start date", TimestampType(), False),\n    StructField("End date", TimestampType(), False),\n    StructField("Start station number", StringType(), False),\n    StructField("Start station", StringType(), False),\n    StructField("End station number", StringType(), False),\n    StructField("End station", StringType(), False),\n    StructField("Bike number", StringType(), False),\n    StructField("Member type", StringType(), False),\n    StructField("timestamp", TimestampType(), False)\n])
```

Priprema podataka za deserijalizaciju is JSON formata.

```
# Parse the "value" field as JSON format and cast the columns to the appropriate data types\nparsed_values = df.select("timestamp", from_json(col("value").cast("string"), schema).alias("parsed_values"))
```

Deserijalizacija.

```
durations = parsed_values.selectExpr("timestamp", "parsed_values.Duration AS Duration", "parsed_values[\"Start station\"] as start_station")
```

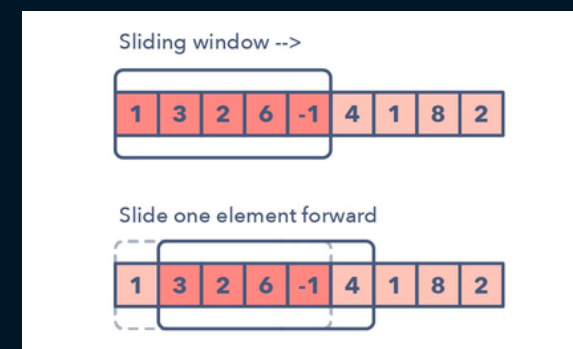
Izdvajanje atributa od interesa.

```
windowDuration = "60 seconds" # The length of the window\nslideDuration = "5 seconds" # The sliding interval
```

Kreiranje klizajućeg prozora i grupisanje podataka po polaznim stanicam

```
durationInfo = durations.groupBy(durations.start_station, window(durations.timestamp, windowDuration, slideDuration)).agg(\n    avg("Duration").alias("avg_duration"),\n    min("Duration").alias("min_duration"),\n    mean("Duration").alias("mean_duration"),\n    max("Duration").alias("max_duration"),\n    count("Start_station").alias("start_station_count"),\n    col("window.start").alias("start_date"),\n    col("window.end").alias("end_date")\n).drop("window")#.dropDuplicates()
```

```
durationInfo.printSchema()
```



```
popular_start_stations = (durations\n    .groupBy(durations.start_station, window(durations.timestamp, windowDuration, windowDuration))\n    .agg(count("*").alias("popularity_count"))\n    .orderBy(desc("popularity_count"))\n    .select(col("start_station"), col("popularity_count"), col("window.start").alias("start_date"), col("window.end").alias("end_date"))\n)
```

Pronalazanje top N najpopularnijih polaznih stanic

```
top_N_start_stations = popular_start_stations.limit(N)\ntop_N_start_stations.printSchema();
```

Upis podataka u bazu

```
query = (durationInfo
        #.withWatermark("timestamp", "1 minute")
        .writeStream
        .outputMode("update")
        .queryName("DeesriptiveAnalysis")
        #.format("console")
        #.trigger(processingTime="5 seconds")
        #.option("truncate", "false")
        .foreachBatch(writeToCassandra)
        .start()
)
```

```
def writeToCassandra(df, epochId):
    df.write \
        .format("org.apache.spark.sql.cassandra") \
        .options(table="sparkone", keyspace="newkeyspace") \
        .mode("append") \
        .save()
    df.show()
```

```
query1 = [
    top_N_start_stations
    .writeStream
    .outputMode("complete")
    .queryName("top_N_start_stations")
    #.format("console")
    .trigger(processingTime="5 seconds")
    #.option("truncate", "true")
    .foreachBatch(writeToCassandra1)
    .start()
]
```

Upiti rade istovremeno

```
query.awaitTermination()
query1.awaitTermination()
```

```
def writeToCassandra1(df, epochId):
    df.write \
        .format("org.apache.spark.sql.cassandra") \
        .options(table="sparktwo", keyspace="newkeyspace") \
        .mode("append") \
        .save()
    df.show()
```

Prikaz rezultata

start_station	start_date	end_date	avg_duration	max_duration	mean_duration	min_duration	start_station_count
Pennsylvania & Minnesota Ave SE	2023-02-23 21:03:40.000000+0000	2023-02-23 21:03:50.000000+0000	546	546	546	546	1
Pennsylvania & Minnesota Ave SE	2023-02-23 21:03:45.000000+0000	2023-02-23 21:03:55.000000+0000	546	546	546	546	1
Pennsylvania & Minnesota Ave SE	2023-02-23 21:30:25.000000+0000	2023-02-23 21:30:35.000000+0000	1334	1334	1334	1334	1
Pennsylvania & Minnesota Ave SE	2023-02-23 21:30:30.000000+0000	2023-02-23 21:30:40.000000+0000	1334	1334	1334	1334	1
New York Ave & Hecht Ave NE	2023-02-23 21:32:30.000000+0000	2023-02-23 21:32:40.000000+0000	799	799	799	799	1
Iwo Jima Memorial/N Meade & 14th St N	2023-02-23 20:36:30.000000+0000	2023-02-23 20:36:40.000000+0000	5322	5325	5322	5319	2
Iwo Jima Memorial/N Meade & 14th St N	2023-02-23 20:36:50.000000+0000	2023-02-23 20:37:00.000000+0000	1759.5	1791	1759.5	1728	2
Iwo Jima Memorial/N Meade & 14th St N	2023-02-23 20:36:55.000000+0000	2023-02-23 20:37:05.000000+0000	1759.5	1791	1759.5	1728	2
Iwo Jima Memorial/N Meade & 14th St N	2023-02-23 20:44:50.000000+0000	2023-02-23 20:45:00.000000+0000	1408	1409	1408	1407	2
Iwo Jima Memorial/N Meade & 14th St N	2023-02-23 20:44:55.000000+0000	2023-02-23 20:45:05.000000+0000	1408	1409	1408	1407	2
Iwo Jima Memorial/N Meade & 14th St N	2023-02-23 20:45:10.000000+0000	2023-02-23 20:45:20.000000+0000	1166.5	1177	1166.5	1156	2
















2023-02-24 13:50:05	+-----+-----+-----+-----+							
2023-02-24 13:50:05		start_station popularity_count		start_date		end_date		
2023-02-24 13:50:05	+-----+-----+-----+-----+							
2023-02-24 13:50:05		8th & O St NW		3 2023-02-24 12:48:40 2023-02-24 12:48:50				
2023-02-24 13:50:05		Neal St & Trinida...		2 2023-02-24 12:48:50 2023-02-24 12:49:00				
2023-02-24 13:50:05		16th & Harvard St NW		2 2023-02-24 12:48:10 2023-02-24 12:48:20				
2023-02-24 13:50:05		11th & S St NW		2 2023-02-24 12:47:30 2023-02-24 12:47:40				
2023-02-24 13:50:05		4th & C St SW		2 2023-02-24 12:47:40 2023-02-24 12:47:50				
2023-02-24 13:50:05		Massachusetts Ave...		2 2023-02-24 12:48:10 2023-02-24 12:48:20				
2023-02-24 13:50:05		13th & H St NE		2 2023-02-24 12:48:20 2023-02-24 12:48:30				
2023-02-24 13:50:05		M St & Pennsylvan...		2 2023-02-24 12:49:10 2023-02-24 12:49:20				
2023-02-24 13:50:05		Columbia Rd & Bel...		2 2023-02-24 12:48:30 2023-02-24 12:48:40				
2023-02-24 13:50:05		19th St & Constit...		2 2023-02-24 12:47:50 2023-02-24 12:48:00				
2023-02-24 13:50:05	+-----+-----+-----+-----+							
2023-02-24 13:50:05								

start_station	start_date	end_date	popularity_count
Iwo Jima Memorial/N Meade & 14th St N	2023-02-24 12:39:30.000000+0000	2023-02-24 12:39:40.000000+0000	2
Iwo Jima Memorial/N Meade & 14th St N	2023-02-24 12:47:30.000000+0000	2023-02-24 12:47:40.000000+0000	2
Iwo Jima Memorial/N Meade & 14th St N	2023-02-24 12:47:50.000000+0000	2023-02-24 12:48:00.000000+0000	2
11th & Kenyon St NW	2023-02-24 12:41:30.000000+0000	2023-02-24 12:41:40.000000+0000	2
Smithsonian-National Mall / Jefferson Dr & 12th St SW	2023-02-24 12:41:30.000000+0000	2023-02-24 12:41:40.000000+0000	2
Smithsonian-National Mall / Jefferson Dr & 12th St SW	2023-02-24 12:44:20.000000+0000	2023-02-24 12:44:30.000000+0000	4
Smithsonian-National Mall / Jefferson Dr & 12th St SW	2023-02-24 12:50:00.000000+0000	2023-02-24 12:50:10.000000+0000	2

Flink aplikacija

- Java 1.8

```
networks:
  bde:
    external: true
services:
  jobmanager:
    image: flink:latest
    expose:
      - "6123"
    ports:
      - "8081:8081"
    command: jobmanager
    environment:
      - JOB_MANAGER_RPC_ADDRESS=jobmanager
    networks:
      - bde
  taskmanager:
    image: flink:latest
    expose:
      - "6121"
      - "6122"
    depends_on:
      - jobmanager
    command: taskmanager
    links:
      - "jobmanager:jobmanager"
    environment:
      - JOB_MANAGER_RPC_ADDRESS=jobmanager
    networks:
      - bde
```

<input type="checkbox"/>	 flink-java	-	Running (2/2)	  
<input type="checkbox"/>	 taskmanager-1 a37652cb0929 	flink:latest	Running	26 seconds ago   
<input type="checkbox"/>	 jobmanager-1 37071988352b 	flink:latest	Running 8081:8081 	29 seconds ago   

localhost:8081/#/overview

Apache Flink Dashboard

- Overview
- Jobs
- Running Jobs
- Completed Jobs
- Task Managers
- Job Manager
- Submit New Job

Version: 1.16.1 | Commit: DeadD0d0 @ 1970-01-01T01:00:00+01:00 | Message: 0

Available Task Slots

1

Total Task Slots 1 | Task Managers 1

Running Jobs

0

Finished 0 | Canceled 0 | Failed 0

Running Job List

Job Name	Start Time	Duration	End Time	Tasks	Status
No Data					

Completed Job List

Job Name	Start Time	Duration	End Time	Tasks	Status
No Data					

Postavke

```
final StreamExecutionEnvironment env = StreamExecutionEnvironment.getExecutionEnvironment();
String inputTopic = "bikes-spark";
//String server = "localhost:9091"; /*app na lokalnu*/
String server = "kafka:9092"; /*app na kontejneru*/
```

Kreiranje klase koja omogućava pokretanje flin streaming programa.

Kafka topic sa koga se citaju podaci.

Kafka server.

```
KafkaSource<String> source = KafkaSource.<String>builder()
    .setBootstrapServers(server)
    .setTopics(inputTopic)
    .setGroupId("my-group")
    .setStartingOffsets(OffsetsInitializer.earliest())
    .setValueOnlyDeserializer(new SimpleStringSchema())
    .build();
```

Kreira se kafka izvor i implementiraju postavke.

```
DataStream<String> text = env.fromSource(source, WatermarkStrategy.nowatermarks(), "sourceName: Kafka Source");
//text.print();
DataStream<BikesTrip> tripDataStream = ConvertStreamFromJsonToBikesTripType(text);
//tripDataStream.print();
```

Kreiranje flink DataStream objekta od kafka izvora.

Mapiranje JSON podataka u BikesTrip format.

```
public static void main(String[] args) throws Exception {

    int n;
    List<String> locations;
    if (args.length < 1) {
        System.out.println("Please provide at least one argument for n");
        return;
    }
    else{
        // Parse program arguments
        n = Integer.parseInt(args[0]);
        locations = new ArrayList<>(Arrays.asList(args).subList(1, args.length));
    }

}
```

Ucitavanje parametara glavne funkcije.

```
public static DataStream<BikesTrip> ConvertStreamFromJsonToBikesTripType(DataStream<String> jsonStream) {
    return jsonStream.map(kafkaMessage -> {
        try {
            JsonNode jsonNode = new ObjectMapper().readValue(kafkaMessage, JsonNode.class);
            return new BikesTrip(jsonNode.get("Duration").asInt(),
                jsonNode.get("Start date").asText(),
                jsonNode.get("End date").asText(),
                jsonNode.get("Start station number").asInt(),
                jsonNode.get("Start station").asText(),
                jsonNode.get("End station number").asInt(),
                jsonNode.get("End station").asText(),
                jsonNode.get("Bike number").asText(),
                jsonNode.get("Member type").asText());
        } catch (Exception e) {
            return null;
        }
    }).filter(Objects::nonNull).forward();
}
```

Iz json-a u BikesTrip.

Funkcionalnost

```
DataStream<BikesTrip> newTripData1 = tripDataStream.rebalance();

SingleOutputStreamOperator<Tuple6<String, Double, Double, Double, Double, Integer>> AggregateStream = newTripData1
    .keyBy(BikesTrip::getStart_station_number) KeyedStream<BikesTrip, Integer>
    .window(SlidingProcessingTimeWindows.of(Time.seconds(20), Time.seconds(10))) WindowedStream<BikesTrip, Integer, TimeWindow>
    .aggregate(new CustomAggregate(locations, threshold: 50));

//.aggregate(new CustomAggregate(locations, 50));

AggregateStream.print();

DataStream<BikesTrip> newTripData = tripDataStream.rebalance();
//int n = 5; // number of most popular start stations to show
SingleOutputStreamOperator<Tuple4<String, String, List<String>, List<Integer>>> popularStationsStream = newTripData
    .windowAll(SlidingProcessingTimeWindows.of(Time.seconds(20), Time.seconds(10)))
    .process(new TopNMostPopular(n));

//popularStationsStream.print();

CassandraService cassandraService = new CassandraService();
cassandraService.sinkToCassandraDB(AggregateStream, table: "tabela", popularStationsStream, table2: "popular_table");

env.execute();
```

Koriscenje funkcije agregacije u kombinaciji sa tehnikom "sliding window" kako bi se dobile informacije o nesledjenim podacima

Pronalazenje top N najpopularnijih polaznih stanica u okviru prozora koriscenjem funkcije process.

Pozivanje cassandra servisa za upis podataka u bazu.

```
public class CustomAggregate implements AggregateFunction<BikesTrip, Tuple5<String, Double, Integer, Double, Integer>, Tuple6<String, Double, Double, Double, Double, Integer>>
public class TopNMostPopular extends ProcessAllWindowFunction<BikesTrip, Tuple4<String, String, List<String>, List<Integer>>, TimeWindow>
```

Upis podataka u bazu

```
public final void sinkToCassandraDB(SingleOutputStreamOperator
```


Pokretanje aplikacije

Aplikacija može biti pokrenuta na lokalu i na kontejnerima.

Prvo je potrebno proslediti parametre glavne funkcije, što su u ovom slučaju broj N i polazna stanice od interesa.

Da bi aplikacija bila pokrenuta na kontejnerima, potrebno je prvo izvršiti komandu **mvn clean package** kako bi kreirali **.jar** fajl aplikacije koji se postavlja na web UI jobmanager-a.

```
C:\Users\Nikola Petrovic\Desktop\bigdata\kafka\docker-kafka-spark-flink-py-java\flink-java>mvn clean package
```

The screenshot displays the Apache Flink Dashboard interface. The sidebar on the left contains navigation links: Overview, Jobs, Running Jobs, Completed Jobs, Task Managers, Job Manager, and Submit New Job. The main area is titled 'Uploaded Jars' and shows a table with columns: Name, Upload Time, and Entry Class. A red arrow labeled 'Step 1.' points to the 'Submit New Job' link in the sidebar. Another red arrow labeled 'Step 2.' points to the '+ Add New' button. A third red arrow labeled 'Step 3.' points to the 'projekat.DataStreamJob' entry in the table. A fourth red arrow labeled 'Step 4.' points to the '5 "Lincoln Memorial" "15th & P St NW"' configuration field. A fifth red arrow labeled 'Step 5.' points to the 'Submit' button. Below the 'Uploaded Jars' section, there is a 'Flink Streaming Job' section showing the job status as 'RUNNING' with 3 tasks. The job ID is 434fdb49fd7a1df19045b0597f7c3f19 and the start time is 2023-02-24 14:34:12. The duration is 46s. The bottom navigation bar includes links for Overview, Exceptions, TimeLine, Checkpoints, and Configuration.

Apache Flink Dashboard

Version: 1.16.1 | Commit: DeadD0d0 @ 1970-01-01T01:00:00+01:00 | Message:

Overview

Jobs

Running Jobs

Completed Jobs

Task Managers

Job Manager

Submit New Job

Uploaded Jars

+ Add New

Name	Upload Time	Entry Class
flink-java-1.0-SNAPSHOT.jar	2023-02-24, 14:32:01	projekat.DataStreamJob

projekat.DataStreamJob

5 "Lincoln Memorial" "15th & P St NW"

Allow Non Restored State

Parallelism

Savepoint Path

Show Plan

Submit

Step 1.

Step 2.

Step 3.

Step 4.

Step 5.

Flink Streaming Job

Cancel Job

Job ID	Job State	Actions
434fdb49fd7a1df19045b0597f7c3f19	RUNNING 3	Job Manager Log

Overview

Exceptions

TimeLine

Checkpoints

Configuration

Rezultati

cqlsh> select* from newkeyspace.tabela;

start_station	avg_duration	min_duration	max_duration	total_duration	count
Lincoln Memorial	2248.5	2206	2322	8994	4
15th & P St NW	201	201	201	201	1
15th & P St NW	730.38788	137	5891	1.2051e+05	165
Lincoln Memorial	1833.5	1828	1839	3667	2
15th & P St NW	786	786	786	786	1
15th & P St NW	419	419	419	419	1
15th & P St NW	731.5283	137	5891	1.5508e+05	212
Lincoln Memorial	1984.8382	362	10137	8.8325e+05	445
Lincoln Memorial	1966.71963	362	10137	1.0522e+06	535
Lincoln Memorial	1714	1714	1714	1714	1
Lincoln Memorial	488	488	488	488	1
Lincoln Memorial	1200.5	811	1590	2401	2
15th & P St NW	665	665	665	665	1
Lincoln Memorial	1019.5	587	1452	2039	2
15th & P St NW	729.99398	137	5891	1.2118e+05	166
Lincoln Memorial	1135	1089	1187	6810	6
Lincoln Memorial	1110	587	1590	4440	4

window_start	window_end	counts	stations
2023-02-24T13:36:40	2023-02-24T13:37	[3, 3, 2, 2, 2]	['Jefferson Dr & 14th St SW', '6th & H St NE', '4th & East Capitol St NE', '17th & G St NW', '5th & K St NW']
2023-02-24T13:37	2023-02-24T13:37:20	[3, 2, 2, 1, 1]	['23rd & E St NW ', 'Massachusetts Ave & Dupont Circle NW', 'North Capitol St & F St NW', '7th & R St NW / Shaw Library', '11th & F St NW']
2023-02-24T13:36:50	2023-02-24T13:37:10	[3, 2, 2, 1, 1]	['23rd & E St NW ', '17th & G St NW', 'Court House Metro / 15th & N Uhle St ', '7th & R St NW / Shaw Library', '12th & L St NW']



**Hvala na
Paznji!**